

# Slate Recommendation Systems: A Reinforcement Learning Problem

## Introduction to Problem Statement:

The PS here deals with building a Recommender System that optimally chooses and places  $C=9$  items on a grid out of a total dataset of  $T=350$  items. This should consider some basic features of the users, expected to be single-use, as well as interaction and correlation amongst the chosen items. We also wish to incorporate features such as seasonality and trends.


## What a Multi Arm Bandit Approach would look like:

A simple bandit approach would treat every slate as one item, i.e. one action set. This would create  $350P9$ , i.e.  $7 \times 10^{22}$  items. Finding enough data to exploit the best actions from this set would be near impossible. This makes it intractable to use on such a dataset. This leads to the need to look at alternate methods of solving this problem

## SlateQ:


A recent approach to the slate recommendation problem is an MDP model called SlateQ. SlateQ looks at building a two-part solution to this problem as follows:

1. A User Choice Model:
  - a. Input:  
Interaction Data from users
  - b. Output:  
The user choice model outputs a relative probability of selecting each item for every grid. For example, this probability would be  $1/10$  for each item before training. (The 10th item being a null item)  
After that, it would then train to calculate relative probability.
  - c. How:  
SlateQ does not explain this model. Here we are looking to ideate and create an alternate model.
  - d. Diagram:



## 2. A Slate Building Model:

- a. Input:  
The relative probabilities of selection, and the rewards associated with each item
- b. Output:  
The best slate out of the  $7 \cdot 10^{22}$  items.
- c. How:  
SlateQ uses mixed integer programming to solve a combinatorial problem in linear time.
- d. Diagram:



## Approach to the problem:

This problem can be approached in two stages:

1. Stage 1: Looking for alternate slate recommendation solutions
  - a. This includes looking at SlateQ and exploring alternatives such as off-policy evaluation, SEO models, and CMAB models.
2. Stage 2: Proceeding with SlateQ and looking at possible User Choice models
  - a. If we proceed with SlateQ, it still leaves uncertainty about how the user choice will be modelled.
  - b. This could include looking at Multinomial Logit models, CTR models or building a model from scratch.

## Stages of solution:

1. Literature Review:
  - a. This is the current step of the process that I am at, which involves reading both about possible alternate models as well as possible options for user choice and understanding their pros and cons
  - b. I have reviewed the papers listed below and am still trying to find evidence that interaction data is modelled at any stage of this process.
  - c. I have found papers that take into account the positioning of the items
2. Database selection:
  - a. We worked initially n setting up the Finn.no database, however, since it could not be fixed up, we moved to the Google RecSim database.
  - b. The Recsim database has now been installed and set up.
3. Proposed Alternative Models:
  - a. Proposed Model Type 1

- i. The first proposal is to consider 350C2 interactions. This would involve learning all 350C2 interactions, then choosing the  $k$  items with the highest correlation magnitude will be factored in, while creating the final slate.
- ii. This reduces complexity by eliminating all interaction terms above 350C3, for which we need significantly more data. It also brings us down from 350C2 to  $350 \times k$  interaction terms to consider.  $k$  could be 1,2,5,10, depending on the data.
- iii. The alternative could be to initially reduce the action space from 350 to around 50 based on the popularity of the singular items (i.e level 1 terms) itself
- b. Proposed Model Type 2
  - i. We could look at non-linearising the action space based on suggestions from Prof.Ravi. This could include building the following.
    1. A neural network-based Contextual Bandit
    2. Relational Boosted Bandits
    3. GP-UCB (Gaussian Process optimizer with upper confidence bounds)

## References:

1. SlateQ:
  - a. Paper: <https://www.ijcai.org/proceedings/2019/0360.pdf>
  - b. Paper 2: <https://arxiv.org/pdf/1905.12767.pdf>
  - c. Implementation Code (Raylib): <https://docs.ray.io/en/latest/rllib/rllib-algorithms.html>
  - d. Source Code (Raylib): [https://docs.ray.io/en/latest/\\_modules/ray/rllib/algorithms/slateq/slateq.html#SlateQConfig](https://docs.ray.io/en/latest/_modules/ray/rllib/algorithms/slateq/slateq.html#SlateQConfig)
  - e. Implementations: <https://github.com/collinprather/SlateQ/blob/master/notebooks>
  - f. Detailed Paper: <https://www.arxiv-vanity.com/papers/1905.12767/#S5>
2. RecSim platform
  - a. Paper: <https://arxiv.org/pdf/1909.04847.pdf>
  - b. Implementation: <https://github.com/google-research/recsim/> (Pasted as text due to format issues)
3. Off-policy evaluation: <https://proceedings.neurips.cc/paper/2017/file/5352696a9ca3397beb79f116f3a33991-Paper.pdf>
4. Combinatorial Multi Arm Bandit: <http://proceedings.mlr.press/v28/chen13a.pdf>
5. Clickthrough Data Optimizer: [https://www.cs.cornell.edu/people/tj/publications/joachims\\_02c.pdf](https://www.cs.cornell.edu/people/tj/publications/joachims_02c.pdf)
6. MNL Model: [https://www.public.asu.edu/~hongminl/mnl\\_attributes\\_price%20v5t\\_final.pdf](https://www.public.asu.edu/~hongminl/mnl_attributes_price%20v5t_final.pdf)
7. Cascading Behaviours:
  - a. Cascading Bandits: <https://arxiv.org/pdf/1502.02763.pdf>
  - b. Multiple Cascading Behaviours: <https://arxiv.org/pdf/1809.08161.pdf>
8. General Survey - RL in Recommender Systems: <https://arxiv.org/pdf/2109.10665.pdf>
9. Gaussian Process Optimization in the Bandit Set - <https://arxiv.org/pdf/0912.3995.pdf>
10. Neural Bandits
  - a. Neural Contextual Bandits: <https://arxiv.org/pdf/2107.03144.pdf>
  - b. Neural Contextual Bandits with UCB: <https://arxiv.org/pdf/1911.04462.pdf>
11. RecSim NG: <https://arxiv.org/pdf/2103.08057.pdf>
12. Cluster Bandits: <https://github.com/google-research/recsim/tree/master/recsim> /colab (Pasted as text due to format issues)

